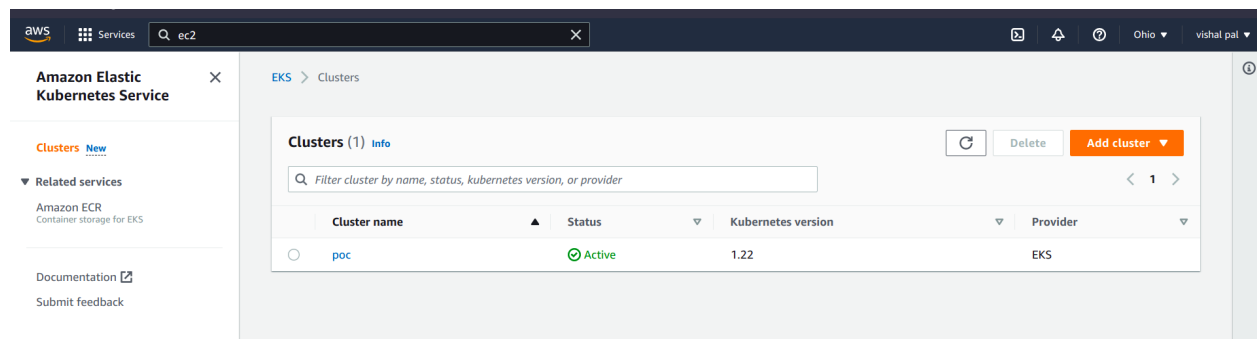


USING SPOT INSTANCE FOR DEPLOYMENT WITH NO DOWNTIME DURING RECLAMATION OF SPOT NODE USING NODE TERMINATION HANDLER.

CREATING EKS CLUSTER

We can refer this document to create a kubernetes cluster.

<https://docs.aws.amazon.com/eks/latest/userguide/create-cluster.html>.

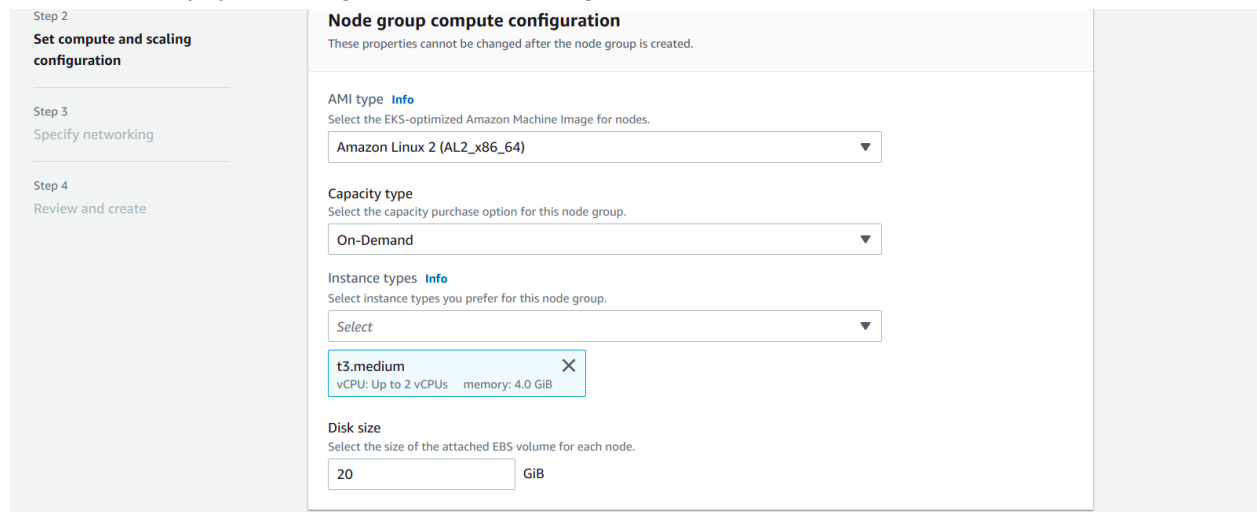


CREATING A ON_DEMAND NODE GROUP

We can refer to this document to create a aws node group.

<https://docs.aws.amazon.com/eks/latest/userguide/create-managed-node-group.html>

In the Capacity type during creation of node group we will select ON_DEMAND



After creation of an on_demand node group we will create a different node group for instance.

We will choose the minimum available node as 1
Maximum available node as 1
Desired nodes as 1
As we are only going to scale SPOT nodes.

CREATING A ON_SPOT NODE GROUP

We can refer to this document to create a spot aws node group.

<https://docs.aws.amazon.com/eks/latest/userguide/create-managed-node-group.html>

During creation of node group we will select capacity type as spot instance

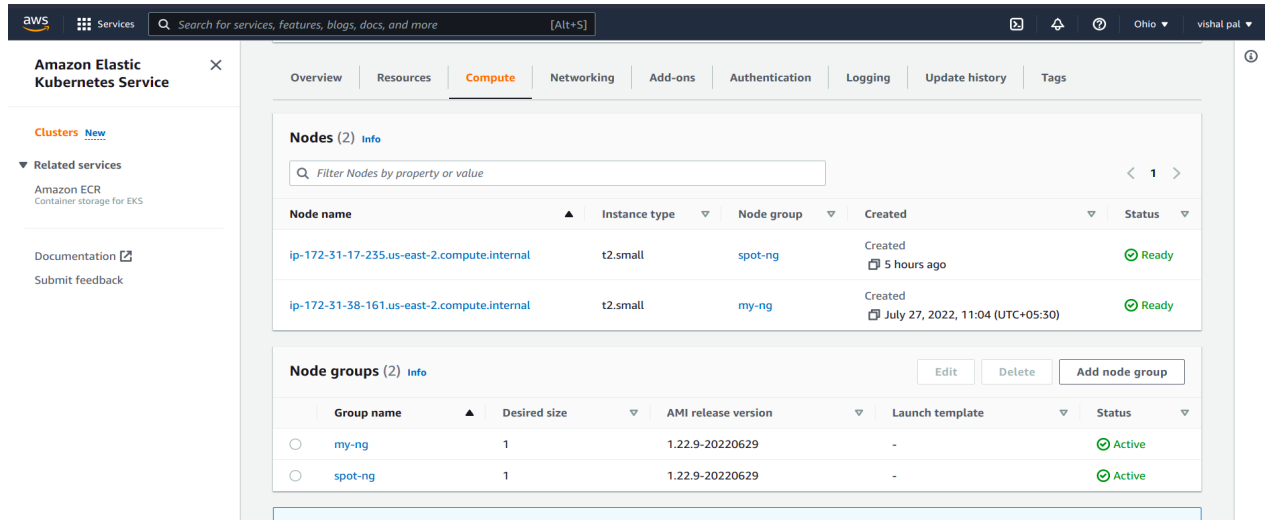
The screenshot shows the 'Node group compute configuration' page in the AWS EKS console. On the left, a sidebar lists four steps: Step 2 (Set compute and scaling configuration), Step 3 (Specify networking), Step 4 (Review and create), and an unselected Step 1. The main content area is titled 'Node group compute configuration' with a warning that properties cannot be changed after creation. It contains four sections: 'AMI type' with a dropdown set to 'Amazon Linux 2 (AL2_x86_64)'; 'Capacity type' with a dropdown set to 'Spot'; 'Instance types' with a dropdown set to 'Select' and a tooltip for 't3.medium' showing 'vCPU: Up to 2 vCPUs' and 'memory: 4.0 GiB'; and 'Disk size' with a dropdown set to '20' and a unit of 'GiB'.

We make sure If we selected Spot for Capacity type, then we recommend specifying multiple instance types to enhance availability.

We will choose the minimum available node as 1
Maximum available node as 4
Desired nodes as 1

Make sure to add a label to this node group.
type = SPOT

After creation of node groups we will have one available node for each node group.



CREATING A DEPLOYMENT AND A SERVICE

We will create a deployment deployment.yaml file and apply it to launch the pods
\$ Kubectl apply -f deployment.yaml

Next we will create a service file for the above deployment
\$ Kubectl apply -f service.yaml

We will serve our deployment using ingress

INSTALLING INGRESS USING HELM

Enter the following command to add the Helm ingress-nginx repo:
\$ helm repo add ingress-nginx <https://kubernetes.github.io/ingress-nginx>

Command to create a namespace for ingress
\$ kubectl create namespace nginx-ingress-sample

Command to deploy ingress resources using helm
\$ helm install my-nginx ingress-nginx/ingress-nginx \
--namespace nginx-ingress-sample \
--set controller.metrics.enabled=true \
--set-string
controller.metrics.service.annotations."prometheus\.io/port"="10254" \
--set-string
controller.metrics.service.annotations."prometheus\.io/scrape"="true"

Now we will create a ingress file for our service

```
unthinkable-lap-0200@PG02PVPL:~/Desktop/aws$ cat myingress
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: minimal-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  #host:
  rules:
  - http:
    paths:
    - path: /pop
      pathType: Prefix
      backend:
        service:
          name: svc-hello
          port:
            number: 8080
```

Now we will apply the file

\$ kubectl apply -f <ingress-file-name>

And now we will get our ingress load balancer url using command:

\$ kubectl get ingress

```
unthinkable-lap-0200@PG02PVPL:~/Desktop/aws$ kubectl get ingress
NAME          CLASS  HOSTS  ADDRESS                                                                 PORTS  AGE
minimal-ingress  nginx  *      a29800e773c144b6cb04cadf8108458f-160032197.us-east-2.elb.amazonaws.com  80     2d9h
```

INSTALLING KUBERNETES METRICS SERVER

<https://docs.aws.amazon.com/eks/latest/userguide/metrics-server.html>

Metrics Server is a scalable, efficient source of container resource metrics for Kubernetes built-in autoscaling pipelines.

Use Case

- CPU/Memory based horizontal autoscaling (learn more about [Horizontal Autoscaling](#))
- Automatically adjusting/suggesting resources needed by containers (learn more about [Vertical Autoscaling](#))

We can use this commands:

\$ kubectl apply -f

<https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml>

To get metrics server deployment

```
$ kubectl get deployment metrics-server -n kube-system
```

INSTALLING HORIZONTAL POD AUTOSCALER

<https://docs.aws.amazon.com/eks/latest/userguide/horizontal-pod-autoscaler.html>

We can directly use this commands :

```
$ kubectl autoscale deployment <deployment-name> --cpu-percent=50 --min=1 --max=10
```

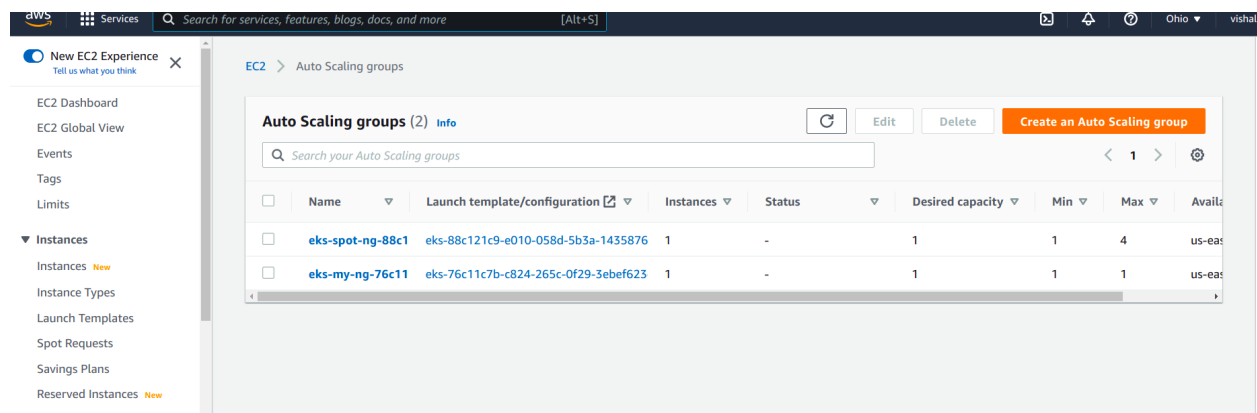
INSTALLING CLUSTER AUTOSCALER ON CLUSTER

The Kubernetes [Cluster Autoscaler](#) automatically adjusts the number of nodes in your cluster when pods fail or are rescheduled onto other nodes. The Cluster Autoscaler is typically installed as a [Deployment](#) in your cluster.

We can refer to this document to install an autoscaler in our eks cluster:

<https://docs.aws.amazon.com/eks/latest/userguide/autoscaling.html>

After installing cluster autoscaler it will create a autoscaling group for both out our node group in AWS autoscaling groups console



We will dynamically scale the spot node group

For creating a dynamic scaling policy

Go to auto scaling groups → Automatic scaling → Dynamic scaling policies → create a Dynamic Autoscaling policy.

We will choose cpu usage to our desired usage after that we want to scale out nodes.
And warmup time for nodes .

Make Sure scale in is enabled.

INSTALLING NODE TERMINATION HANDLER ON CLUSTER

To install node termination handler we will use these command

```
$ helm repo add eks https://aws.github.io/eks-charts
```

```
$ helm install aws-node-termination-handler \
  --namespace kube-system \
  --version 0.15.4 \
  --set nodeSelector.type= SPOT\
  eks/aws-node-termination-handler
```

To manage the Availability of our website we will create two deployments with a single service to handle no downtime during the reclamation of nodes by aws.

To make sure there is at least one pod available on each node we will use selectors to deploy on each node .

We can generate load on our load balancer endpoint using

```
$ kubectl run -i \
  --tty load-generator \
  --rm --image=busybox \
  --restart=Never \
  -- /bin/sh -c "while sleep 0.01; do wget -q -O- <out-end-point>; done"
```

To get HPA we can use:

```
$ kubectl get hpa
```

Thank you.